

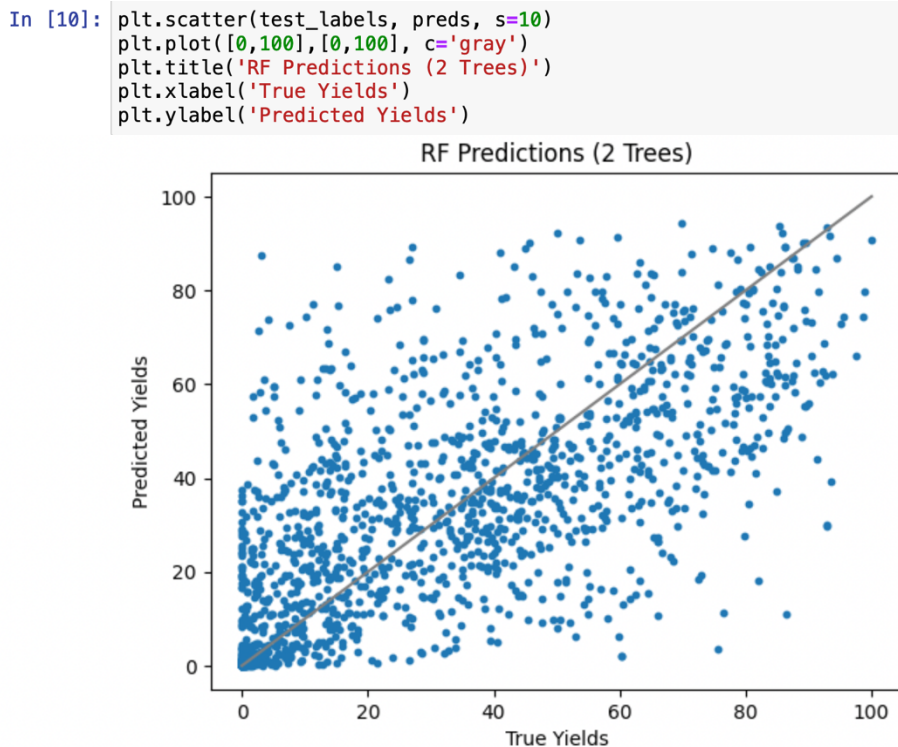
And train our RF.

```
In [7]: rf.fit(train_inputs, train_labels)
Out[7]: RandomForestRegressor(n_estimators=2)
```

We can then predict on our held out test data.

```
In [9]: preds = rf.predict(test_inputs)
print("MAE:", mean_absolute_error(test_labels, preds), "RMSE:", root_mean_squared_error(test_labels,
                                                                                          preds),
      "R2:", r2_score(test_labels, preds))
MAE: 14.285283063831026 RMSE: 20.158044273452525 R2: 0.46532430606747566
```

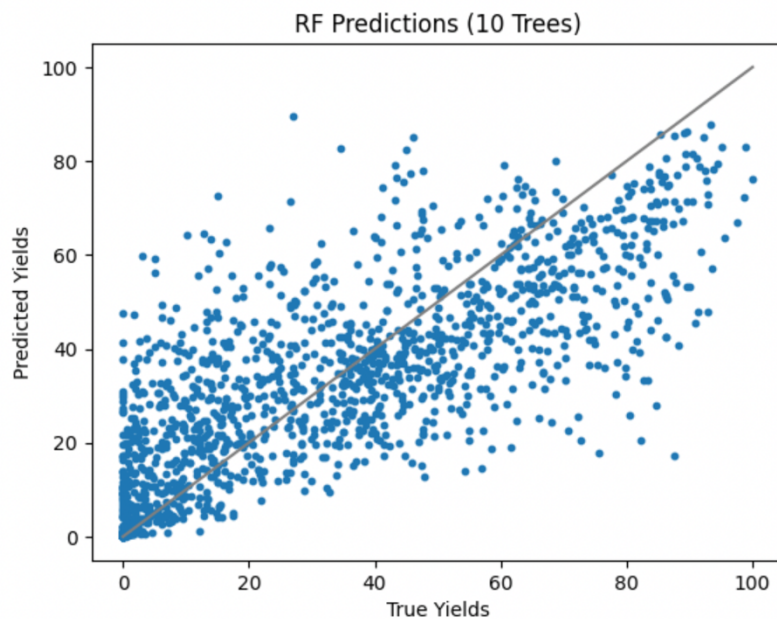
Again we can plot these predictions to get a visual sense of how well we're doing.



What if we increase to 10 trees in our forest? Well, we can define our new RF and re-run the training and predictions. Note that we are using the same test set so as to have an apples to apples comparison.

```
In [11]: rf = RandomForestRegressor(n_estimators=10)
rf.fit(train_inputs, train_labels)
preds = rf.predict(test_inputs)
print("MAE:", mean_absolute_error(test_labels, preds), "RMSE:", root_mean_squared_error(test_labels,
                                                                                          preds),
      "R2:", r2_score(test_labels, preds))
MAE: 12.25451479790481 RMSE: 16.96950840227857 R2: 0.6210933959481579
```

```
In [12]: plt.scatter(test_labels, preds, s=10)
plt.plot([0,100],[0,100], c='gray')
plt.title('RF Predictions (10 Trees)')
plt.xlabel('True Yields')
plt.ylabel('Predicted Yields')
```



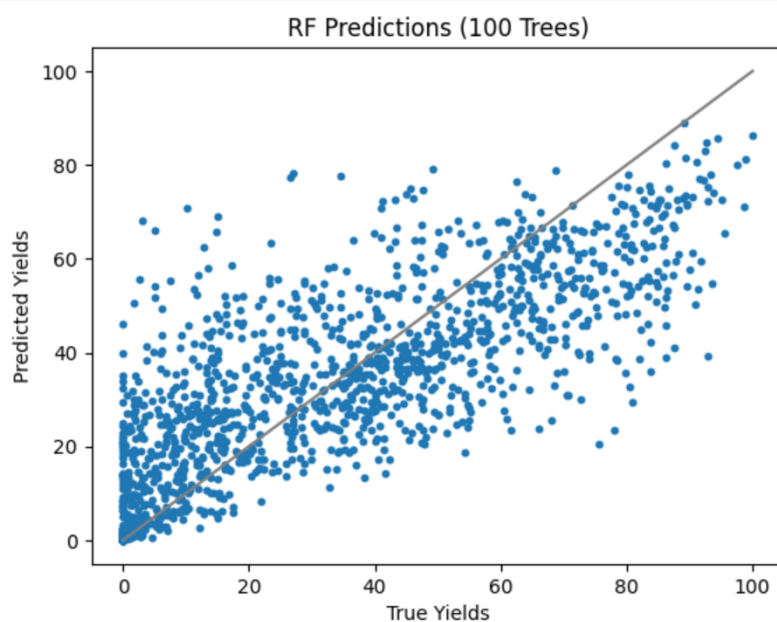
Kind of incredible what just 10 trees can do for you.

Okay, what about 100 trees? This is actually the new default option in scikit.

```
In [13]: rf = RandomForestRegressor(n_estimators=100)
rf.fit(train_inputs, train_labels)
preds = rf.predict(test_inputs)
print("MAE:", mean_absolute_error(test_labels, preds), "RMSE:", root_mean_squared_error(test_labels,
                                                                                          preds),
      "R2:", r2_score(test_labels, preds))
```

MAE: 11.987471793217269 RMSE: 16.354215087248562 R2: 0.6480726148027205

```
In [14]: plt.scatter(test_labels, preds, s=10)
plt.plot([0,100],[0,100], c='gray')
plt.title('RF Predictions (100 Trees)')
plt.xlabel('True Yields')
plt.ylabel('Predicted Yields')
```



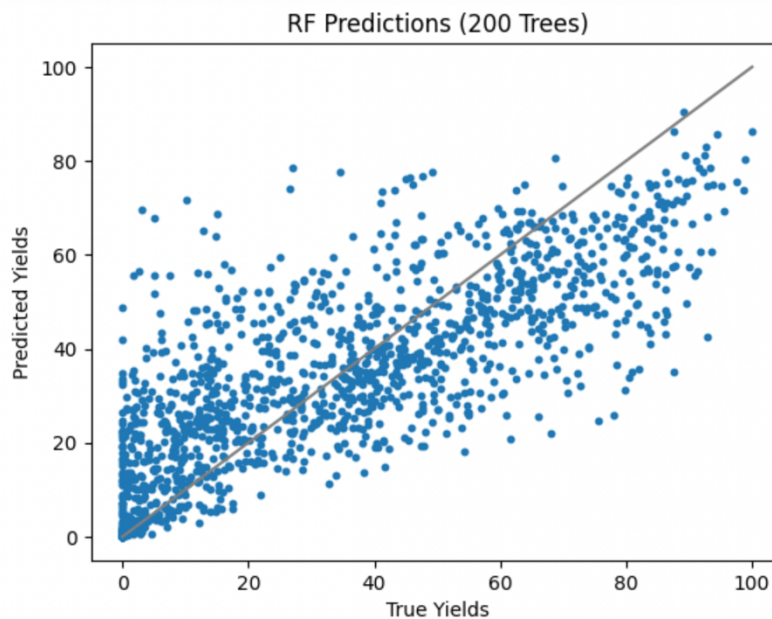
We can see a noticeable improvement, but definitely not as drastic as going from 2 to 10 trees.

How about 200 trees?

```
In [15]: rf = RandomForestRegressor(n_estimators=200)
rf.fit(train_inputs, train_labels)
preds = rf.predict(test_inputs)
print("MAE:", mean_absolute_error(test_labels, preds), "RMSE:", root_mean_squared_error(test_labels,
                                                                                          preds),
      "R2:", r2_score(test_labels, preds))
```

MAE: 11.901714814676033 RMSE: 16.274115675208485 R2: 0.6515115012594233

```
In [16]: plt.scatter(test_labels, preds, s=10)
plt.plot([0,100],[0,100], c='gray')
plt.title('RF Predictions (200 Trees)')
plt.xlabel('True Yields')
plt.ylabel('Predicted Yields')
```



You may start to notice that a RF with this many trees does start to slow you down a bit and to top it off, we really aren't seeing noticeable improvement anymore (although more trees doesn't hurt).

And that's the power of RFs! Really fast, generally pretty accurate, and no preprocessing is needed.

– Happy Coding!